



Concept drift detection in datasets using tree-structures. The total number of trees in the forest model is held constant at **100** for all datasets. Each dataset has a streaming window size commensurate with its total size: *Abalone*(512), *ANN-Thyroid-1v3*(512), *Cardiotocography*(512), *Covtype*(4096), *Electricity*(1024), *KDDCup99*(4096), *Mammography*(4096), *Shuttle*(4096), *Weather*(1024), and *Yeast*(512). The last data window in each dataset usually has much fewer instances and therefore its distribution is very different from the previous window despite there being no data drift. **Therefore, ignore the drift in the last window.** We did not expect *Abalone*, *ANN-Thyroid-1v3*, *Cardiotocography*, *KDDCup99*, *Mammography*, *Shuttle*, and *Yeast* to have much drift in data, and this can also be seen in the plots where most of the windows in the middle of streaming did not result in too many trees being replaced (the numbers in the parenthesis are mostly zero). We also did not expect *Covtype* to have much drift in its default ordering (from its public repository); however, after seeing the plots we realize that there is a definite drift due to which a significant number of trees need to be replaced with each new window of data. *Electricity* and *Weather* are standard streaming datasets with expected concept drift, which shows up in the plots.